

Terminology Extraction

Adam Meyers
New York University



Terminology Extraction

- Given
 - set of documents about a topic (foreground)
 - set of documents about diverse topics (background)
- Find ranked list of terms (words, n-grams, etc.)
 - that are more characteristic of foreground than background
- Uses:
 - terms for previously described tasks, search terms, terms for glossary, terms to track for technology forecasting (predicting technological emergence), etc.



Termolator

- Open Source Terminology Extraction for Chinese, French & English
 - <http://nlp.cs.nyu.edu/termolator/>
- Created under a government contract as part of the Foresight and Understanding from Scientific Exposition (FUSE) and some subsequent support.
- Collaborators at NYU: Zachary Glass, Ralph Grishman, Yifan He, Giancarlo Lee, Shasha Liao, Angus Grieve-Smith, John Ortega, Yuling Gu, Leizhen Shi, Sandra Burlaud, Anand Tyagi and others



What is Terminology?

- Webster's II New Collegiate Dictionary Definition
 - *The vocabulary of technical terms and usages appropriate to a particular field, subject, science, or art.*
- Operational Definitions:
 - Keyword sequences for Information Retrieval (IR)
 - Need not be technical, e.g., *wheat, barley, white mouse*, in genetics
 - Items to define in Technical Glossaries
 - Items to track for Technology Forecasting (TF)
- Noun Terminology:
 - Technical word sequence headed by noun
 - Vast majority of all terminology
 - Non-noun terminology exists, but not included in this research

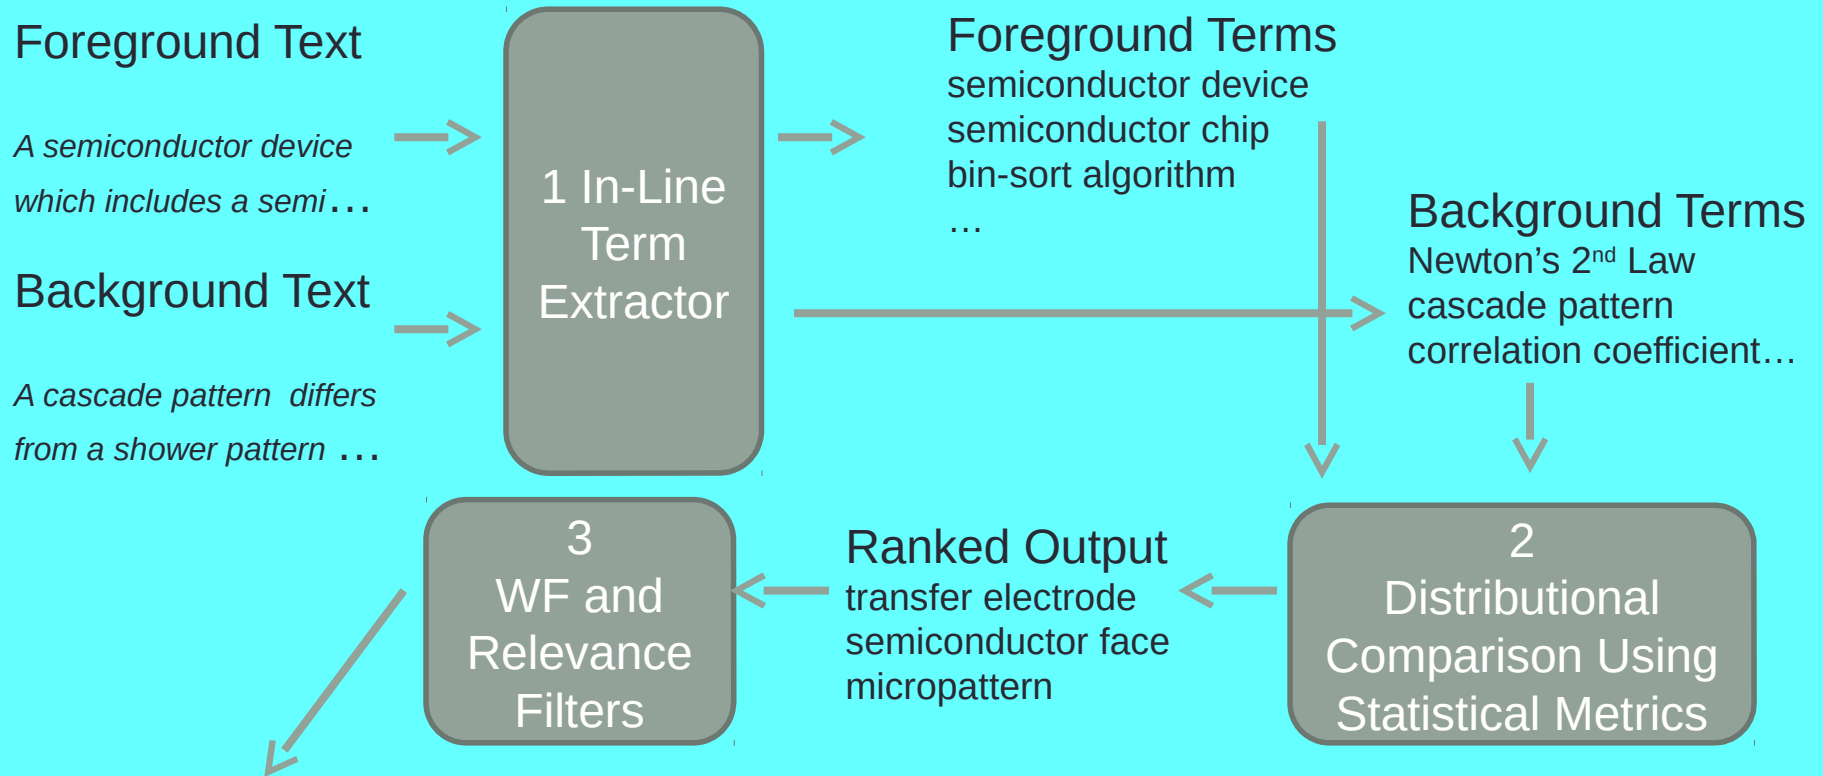


Examples of Terminology

- Juggling: ***cascade pattern, Mills mess, full shower***
- Real Estate: ***balloon mortgage, title search, full shower***
- Computer Science: **hidden markov model, genetic algorithm, top-down search**
- Knitting: ***gobelin stitch, half-treble, corner scallop***
- biology: ***myosin-ii, plasminogen activator, antizyme***



Termolator Work Flow



Final ReRanked Output: transfer electrode; semiconductor face; micropattern, through-connection, wavelength conversion chip, ...



The Termolator: 2 Main Subsystems

- **In-Line Term System:** Finds instances of terms (tokens)
 - Finds noun/adjective sequences that obey constraints
 - Identifies term tokens, instances of terms in sentences
 - 500 term tokens occur in document X
 - 50 are instances of ***biotrophic effector models***
 - Limited previous work in this area
- **Distributional Term System:** Finds term types
 - Counts instances of term types
 - 30 term types occur in document X
 - ***biotrophic effector models*** occurred 500 times
 - Ranks term types by characteristic-ness to a particular topic
 - Top N term types are kept, the rest are discarded
 - Uses metrics similar to TF-IDF discussed in previous slides



Our In-line Term Extraction System

- Manual Rule Based “Chunker”
 - Identifies sequences of nouns and adjectives
 - part of speech tagger output
 - dictionaries
 - Technical words identified:
 - Out-of-Vocabulary (OOV) words – words not in dictionaries
 - ***semiconductor, biotrophic, gobelin***
 - Technical Adjectives – based on endings (-ic,-cal,-ous) and dictionaries
 - ***algebraic, amphibious, umbilical***
 - Nominalizations – based on endings (tion, etc.) and NOMLEX dictionary
 - ***conduction, vulcanization, accelerator***
- Well-formedness filter
 - eliminates ill-formed (too short, bad characters, etc.)
 - eliminates terms without OOV or technical words
 - eliminates words detected to be names of people or places



Example Identification of Technical Noun Adjective Sequences

- A **semiconductor device** which includes: a **semiconductor chip** bonded to a **surface** of a solid **device**; and a **stiffener** surrounding the **periphery** of the **semiconductor chip**.

A_{DET/O} semiconductor_{O-NOUN/B} device_{NOUN/I} which_{OTHER/O} includes_{OTHER/O} a_{DET/O}
semiconductor_{O-NOUN/B} chip_{NOUN/I} bonded_{VERB/O} to_{PREP/O} a_{DET/O} surface_{NOUN/B} of_{PREP/O}
a_{DET/O} solid_{ADJ/O} device_{NOUN/B} ;_{OTHER/O} and_{OTHER/O} a_{DET/O} stiffener_{NOUN/B} surrounding_{VERB/}
O the_{DET/O} periphery_{NOUN/B} of_{PREP/O} the_{DET/O}
semiconductor_{O-OUN/B} chip_{NOUN/I} •_{OTHER/O}

Rules group yellow words (below) together resulting in blue sequences (above).



Filters Remove Unlikely Candidate Terms

- Accepts Terms which contain an Out-of-vocabulary (OOV) word
 - semiconductor/O-NOUN device
 - semiconductor/O-NOUN chip (2 instances)
- Accepts Terms containing technical adjectives or nominalizations
 - thermal/TECH-ADJ stress
 - fabrication/NOM process
- Rejects Terms because they contain no technical words
 - *surface*
 - *device*
 - *stiffener*
 - *periphery*
- Other Non-Terms removed for other reasons
 - T
 - 212-345-8888
 - No.
 - New York



Supplementary patterns for identifying Terms

- Arguments of Abbreviation relations
 - Not organizations or places
 - Aligns words before parentheses with word in parentheses
 - *already been chewed (ABC)*
 - XML (Extensible Markup Language)
 - *third variable loop (V3)*
 - ***D. melanogaster gene Muscle LIM protein at 84B (abbreviated as Mlp84B)***
 - *Schwartz and Hearst (2003)*
- Terms Matching Regexp Patterns
 - Gene Sequences: **AACAAGGTGGCGCAGTT**
 - Chemical Formulas: **Ag2CrO4**



Evaluation of Inline Term System

- 2 Annotators Manually Annotated Inline terms in 3 documents
- Adjudicated the Results
- Scored annotators against adjudicated annotation
- Scored system against adjudicated annotation
- Compared annotator vs system performance



Annotation

- Setup
 - 2 annotators annotated the same three documents
 - Annotator 2 Adjudicated
 - Annotator 1's score against Adjudicated results may be a good Upper Bound for evaluating the Automatic System (assumes the adjudication is biased in favor of Annotator 2).
- Defining Inline Term for Annotator
 - Single or multi-word nominal expression specific to technical discipline
 - It can be conventionalized by defining or abbreviating it early in the document and by reusing the term
 - Is term specific to technical discipline, i.e., is it obscure enough?
 - Would a naïve adult (like Homer Simpson) know the term?
 - Is it found in the Juvenile subcorpus of the Corpus of Contemporary American English (<http://corpus.byu.edu/coca/>)?



Corpora and Systems Tested

- Corpora
 - A Speech Recognition Patent (SRC)
 - A Sun Screen Patent (SUP)
 - A Journal Article about a Virus Vaccine (VVA)
- Systems Tested
 - Base 1: assume all noun groups minus determiners are terms
 - use MEMM chunker with Genia (Kim et al 2003) features
 - Base 2: baseline 1 system, but filtered by only keeping those Noun Groups that end with an O-NOUN
 - System without Filter: Chunking system as described, but without the filter
 - Final System
- Matching Criteria
 - Strict Match – The test term and answer key term are the same
 - Sloppy Match – The test term and answer key term overlap in extent.



Inter Annotator Agreement

				Strict				Sloppy		
	Doc	Terms	Matches	Pre	Rec	F	Matches	Pre	Rec	F
Annot 1	SRP	1131	798	70.8%	70.6%	70.7%	1041	92.5%	92.0%	92.2%
	SUP	2166	1809	87.5%	83.5%	85.5%	1992	96.3%	92.0%	94.1%
	VVA	919	713	90.9%	77.6%	83.7%	762	97.2%	82.9%	89.5%
Annot 2	SRP	1131	960	98.4%	84.9%	91.1%	968	99.2%	85.6%	91.9%
	SUP	2166	1999	95.5%	92.3%	93.8%	2062	98.5%	95.2%	96.8%
	VVA	919	838	97.4%	91.2%	94.2%	855	99.4%	93.0%	96.1%

Annotator 1 scores may be upper bounds for system results



Baseline Systems

				Strict				Sloppy		
	Doc	Terms	Matches	Pre	Rec	F	Matches	Pre	Rec	F
Base 1	SRP	1131	602	24.3%	53.2%	33.4%	968	44.2%	96.8%	60.7%
	SUP	2166	1367	36.5%	63.1%	46.2%	1897	50.6%	87.6%	64.2%
	VVA	919	576	28.5%	62.7%	39.2%	887	44.0%	96.5%	60.4%
Base 2	SRP	1131	66	24.9%	5.8%	9.5%	151	57.0%	13.4%	21.6%
	SUP	2166	771	52.3%	35.6%	42.4%	1007	68.4%	46.5%	55.3%
	VVA	919	270	45.8%	29.4%	35.8%	392	66.5%	42.6%	51.9%

- Base 1 (all noun groups): results in high recall/low precision
- Base 2 (must end in O-NOUN): too severe a filter.



System Results

				Strict				Sloppy		
	Doc	Terms	Matches	Pre	Rec	F	Matches	Pre	Rec	F
No Filter	SRP	1131	932	39.0%	82.4%	53.0%	1121	46.9%	99.1%	63.7%
	SUP	2166	1475	39.7%	68.1%	50.2%	1962	52.8%	90.6%	66.7%
	VVA	919	629	27.8%	68.4%	39.5%	900	39.8%	97.9%	56.6%
Final System	SRP	1131	669	69.0%	59.2%	63.7%	802	82.8%	70.9%	76.4%
	SUP	2166	1193	64.7%	55.1%	59.5%	1526	82.8%	70.5%	76.1%
	VVA	919	581	62.1%	63.2%	62.7%	722	77.2%	78.6%	77.9%

Final System gets the highest F-score



Distributional Term System

- Find In-line Terms for Foreground Corpus (or sample)
- Find In-line Terms for Background Corpus (or sample)
- Count instances of the same lemma as instances of the same term
 - singular/plural, -ing endings (stemming)
 - *speech recognizers* → *speech recognizer*
 - abbreviation/full-form
 - *html* → *hypertext markup language*
 - Noun mod alternations:
 - *Recognition of Speech* → *Speech Recognition*
- Rank by Statistical Metrics similar to TF-IDF
 - finds terms more characteristic of foreground than background
- Rerank terms using Relevance Metric, based on a Yahoo Websearch
- Take Top N terms (e.g., N = 5000)



Statistical Metrics for Ranking Terms

- A linear combination of 3 Measures comparing the distribution of terms in the foreground (For) vs background (Bac)

- Term Frequency Inverse Document Frequency (TFIDF)

$$TFIDF(t) = \frac{FreqFor(t)}{FreqBac(t)} \times \log\left(\frac{NumBacDocuments}{NumBacDocsContain(t)}\right)$$

- Document Relevance Document Consensus (DRDC)

- Navigli and Velardi (2004)

$$DRDC(t) = \frac{FreqFor(t)}{FreqBac(t)} \times \sum_{d \in Foreground} \frac{freqBac(t, d)}{freqFor(t)} \times \log\left(\frac{freqFor(t)}{freqBac(t, d)}\right)$$

- Doc Relevance (1st factor) favors representative terms (like TFIDF)

- Doc Consensus (2nd factor) favors terms found in many documents

- Kullback Leibler Divergence (KLD)

- Cover and Thomas (1991), Hisamitsu, et. al. (1999)

$$KLD(t) = (\log(freqFor(t)) - \log(freqBac(t))) \times freqFor(t)$$

- Compares Probability of term occurs in Foreground vs. Background

-



Filters on Distributional Output

- 2 Filters that can be applied to our system or output of other term generation systems
 - In FUSE, they were applied to MITRE and BBN output
- Both scores are between 0 and 1, they are combined by multiplication
- Well-Formedness Filter
 - Many of the constraints are built into our chunker
 - Most terms have a score of 1
 - However, component of distributional System adds some common substrings of terms to output, some of which are ill-formed
- Relevance Filter
 - We use a Yahoo search result and heuristics to score terms more highly if they are used in articles or patents



Well-Formedness Filter

- A term is well-formed if it is:
 - An abbreviation
 - A set of words that is abbreviated somewhere in the corpus
 - A single out of vocabulary word
 - Matches a regular expression that finds chemical names, DNA sequences or paths (urls, bio paths, etc.) – although URLs can be documents, rather than terms.
- A term is also well formed if it obeys noun group rules (a sequence of adjectives and nouns ending in a noun) AND it contains at least one out-of-vocabulary word, nominalization or technical adjective
- The degree of ill-formedness is not so important as scores below 1 rarely apply to accepted terms. (Sometimes favors terms with OOV words over terms with other technical words and no OOVs)
- This filter is more important when applied to term lists not created by The Termolator (Mitre and BBN term lists in FUSE)



Relevance Filter

- Run on each term below some cutoff (typically 30K)
 - Time consuming (about .75 seconds per term)
- Yahoo search (Bing) for exact match of term
- Relevance = H^2T
 - $H = 0$ to 1 score based on number of hits
 - $$\frac{\min(\log_{10}(\text{numberHits}), 10)}{10}$$
 - Minimized for non-hits (0 hits counts as 500 hits)
 - T = Percent of top 10 hits that are articles or patents
 - Based on key word search in title, url & summary
 - Key words = {patent, article, proceedings, journal, dissertation, abstract, ...}



Evaluation of Termolator Output

- Foreground Corpus: 2500 patents about optical systems
 - US Patent codes: 250, 349, 356, 359, 362, 385, 398 and 399
- Background Corpus: 2500 randomly selected patents
- Years: 1997-2007
- Took the top 30K out of 219K terms and reranked using:
 - Percentile X Well-Formedness X Relevance
- Manually evaluated 100 terms sampled from top 5000 terms
- A term was judged correct if
 - valid keyword
 - not missing crucial modifier
 - did not contain any spurious word.
- **The system achieved 86% Precision**
- Recall difficult to measure, but also produces more high-quality terms
- Competitive with other systems (main innovation is: inline terms)



Evaluation Details

- Sample Correct (sampled from the first 5000):
 - *stimulable phosphor, ion beam profile, x-ray receiver, wavelength-variable, quadrupole lens, proximity correction, dfb laser, asymmetric stress, panoramagram, single-mode optical fiber, total reflection plane, photosensitive epoxy resin*
- Sample Incorrect
 - *irradiation time t*
 - A variable, not a term (without *t*, it would be a term)
 - *evolution*
 - This word has entered the common vocabulary
 - *crystal adjacent*
 - This word sequence includes two words at a constituent boundary
 - a noun phrase followed by a modifying adjective phrase, e.g.,
 - *[[a liquid crystal] [adjacent to the lower alignment layer]]*



Informal Observations about Recall

- Recall or coverage is difficult to measure without an exhaustive amount of human annotation
- The distributional system gets roughly the same precision for Noun Group input as Inline Term Group Input for the top N terms, where N is a small number
- Using Inline Terms as input, we generate many more terms with high scores and thus seem to improve Recall by a large amount (at least a factor of 2)
 - But this is hard to measure
- Rationale: Garbage In → Garbage Out
 - High F-scores for inline terms (vs NGs or N-grams)
 - Higher Quality terms are being ranked and so the high-ranked items are more likely to be correct

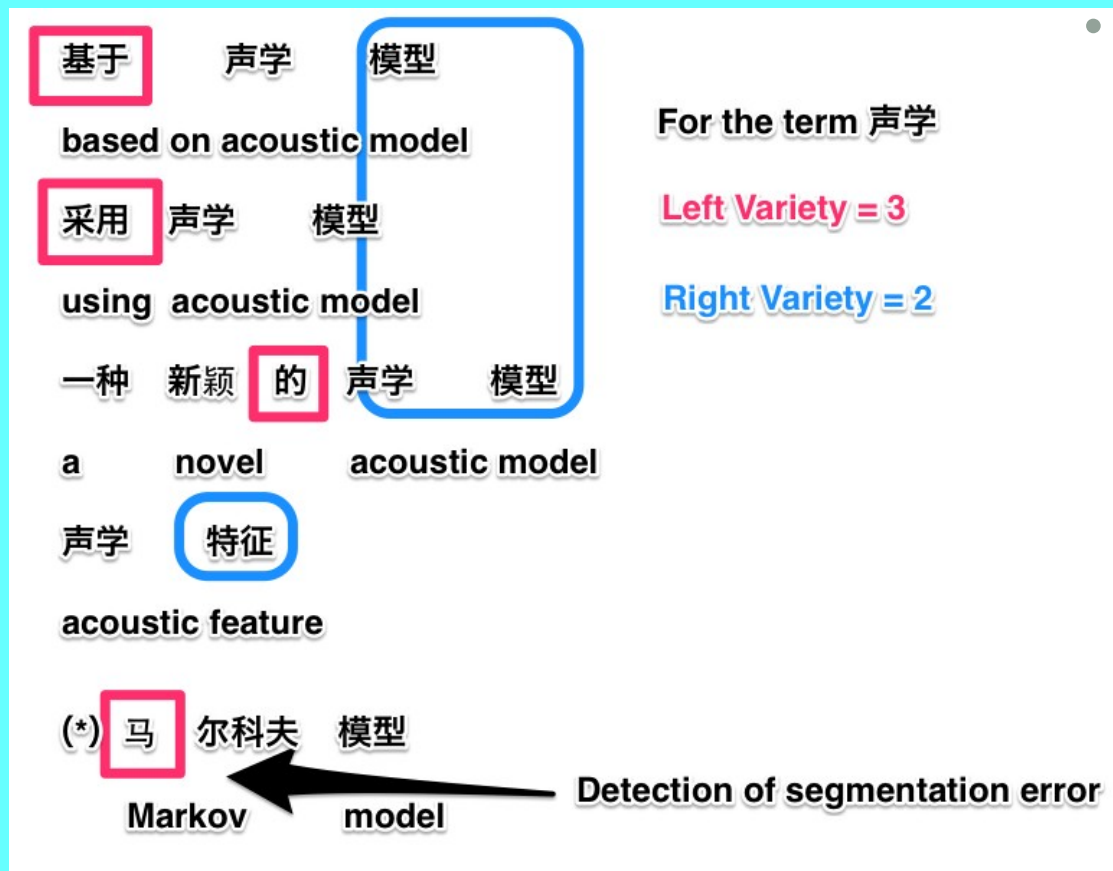


The Termolator for Chinese 🕶️

- Work by Yifan He
- Distributional System is the Same as English
- Uses Noun Group Chunker for input terms
- Accessor-Variety Filter (Feng et al., 2004)
 - Score Based on the Number of distinct words that appear before and after a particular term type
 - Low Scores indicate unlikely Chinese words
- 1100 terms extracted from 2000 speech recognition patents
 - 78% precision on top 50 terms
 - 85% precision on top 20 terms
- 2018-2019 research by students Y. Gu & L. Shi
 - Unifying components of English and Chinese System
 - Additional Chinese Components and Evaluation (in-progress)



Example of Chinese Term Filtering




Examples for Access Variety based filtering

- 尔科夫模型 (Markov model, with the first Chinese character 马 missing) is probably a boundary error
- [Pic on left] 尔科夫模型 has the same character 马 on its left boundary thus its Left AV=1
- [Pic on left] A correct term 声学 (acoustics) will have Left AV>=3



Open Source Distribution

- Open Source release of The Termolator 
 - NYU's Website: <http://nlp.cs.nyu.edu/termolator/>
 - Github:
 - English: https://github.com/AdamMeyers/The_Termolator
 - Future release planned of multi-lingual system
 - Chinese: <https://github.com/ivanhe/termolator/>
- English system for UTF-8 (including ASCII) & ISO-8859-1
- Tested on Public Domain Texts
 - Google Patents
 - Project Gutenberg
 - Open American National Corpus



Examples from Public Domain Texts

- Gutenberg: Chapters in a Book about knitting vs Other Docs
 - *open-work insertion, fine mesh, transverse stitching, empty scallop*
- Open American National Corpus (OANC) – Biology documents versus random documents
 - *myosin-ii, hsn3, intron, migration defect, sparc-null mice*
- Google Patents: Surgery patents (US Patent Class 606) vs Random Patents:
 - *fluid manifold, dissector arm, pedicle punch, balloon catheter*



Subsequent Work on Termolator

- Character-based language model:
 - Eliminate unlikely paragraphs (-1 stand dev from mean)
 - Eliminates bibliographies, figures, noise
 - Model based on 5-grams of sequences of
 - WHITESPACE
 - DIGIT
 - LETTER
 - PUNCTUATION
 - OTHER
- Merge Chinese and English system (Yuling Gu, Leizhen Shi, Echo Hong, Yuting Wang)
 - Some English components are not possible for Chinese
 - We could not initially improve over standard noun chunker, but Yuting is experimenting with some knowledge-based information.
 - Websearch does not help (too few open source technical documents)
 - Chinese Character-based language model seems OK
 - Some components are Chinese or English specific, but some components (distributional) are shared
- French system – preliminary, but all English components seem possible (Sandra Burland)
- System for getting foreground and background documents from Wikipedia (Sandra Burland, Anand Tyagi)
- **Recent/Current Work on upcoming slides**



Current and Recent Research on Termolator



Termolator

- A. Meyers, Y. He, Z. Glass, J. Ortega, S. Liao, A. Grieve-Smith, R. Grishman and O. Babko-Malaya (2018). “*The Termolator: Terminology Recognition based on Chunking, Statistical and Search-based Scores.*” *Frontiers in Research Metrics and Analytics*
 - <https://www.frontiersin.org/articles/10.3389/frma.2018.00019>
- N. Pham, L. Pham and A. Meyers (2021) “Legal Terminology Extraction with the Termolator”. *NLLP-2021*.
 - <https://aclanthology.org/2021.nllp-1.16/>
- S. Nordquist and A. Meyers (2022). “On Breadth Alone: Improving the Precision of Terminology Extraction Systems on Patent Corpora”. *NLLP-2022*
- Code (a few updates are forthcoming)
 - <http://nlp.cs.nyu.edu/termolator/>
 - https://github.com/AdamMeyers/The_Termolator
 - <https://github.com/ivanhe/termolator/>



Termolator – Current Work

- Several improvements to the English system since paper was written
 - Automatic Glossary for terminology output (evaluation step still required)
 - Automatic Glossary from Wikipedia: with Anand Tyagi
 - Language Model for removing ill-formed input (tables, bibliography)
 - System for getting input corpora from Wikipedia (Sandra Burlaud)
- Integration of Chinese and French Versions (Leizhen Shi, Echo Hong, Yuling Gu, Sandra Burlaud)
- Initial German and Arabic Versions (past class projects)
- Tuning English system to court opinions – Lachlan Pham and Nhi Pham
 - Paper at NLLP 2021: <https://aclanthology.org/2021.nllp-1.16.pdf>
- Work on evaluating the effect of different types of foreground and background corpora
 - Sean Nordquist (To be submitted to Terminology Workshop at LREC 2022)
- Future Work:
 - Chinese System
 - Other languages
 - Evaluation for Glossary



Termolator Text Front End and Glossaries

- Demo of Command Line Tool
 - May not complete running until after class
 - Running System
 - Keyword selection of Foreground from Wikipedia
 - Background selected based on Superclasses found in Wikipedia
- Creates “Glossary”



Glossary Continued

- Glossary takes highest scoring terms that occur in at least 3 documents
- Displays one “entry” for each term
 - Wikipedia First Paragraphs
 - 3 “diverse” sample paragraphs
- Examples:
 - Machine_learning – .edited_term_list and .summary
 - test_summaries
 - Category Theory (bk: all)
 - Printmaking (bk: "Artistic techniques" + "Visual arts media")



1st Paragraphs from Wikipedia

- If the whole string matches a Wikipedia Entry, that first paragraph is used, e.g., the first few matches for “Category Theory”
- If not, but substrings match Wikipedia, the substring 1st paragraphs are used, e.g., “anilox roll” under “Printmaking”



Selecting the Sample Usages

- Create vectors of TF-IDF scores for all the examples paragraphs containing the term.
- Compute a centroid vector – the average of all these vectors.
- Form 3 clusters of vectors using, a simple clustering method:
 - iteratively merge most similar clusters one at a time until there are only 3 clusters left
- From each cluster, select the example that is the least similar to the centroid.



Status

- Problem1: not publishable without formal evaluation
- Problem2: Testable for English and probably French (with minor help), but need help with Chinese version (including debugging)
- Informal Evaluation: probably helps understand terms in a technical domain
- Formal Evaluation Required
 - Human Evaluation?
 - Likert Scores reflecting inclusivity, clarity
 - Number of Errors
 - Automatic Evaluation
 - Comparison with manually created glossaries (BLEU, ROUGE, etc.)
- Problem2:
 - Works for English
 - Testable for French with minor debugging
 - Chinese requires some debugging plus there are some issues with variation between simplified and traditional characters in Wikipedia



Tuning Termolator to the Court Decisions

- Data: Supreme Court Database (SCDB) from Washington University School of Law
 - Spaeth et. al. 2013 <http://supremecourtdatabase.org>
 - Can be downloaded through Python's Textacy library
- Our Research: N Pham, L. Pham and A. Meyers (2021)
- Tuned Termolator to work on Supreme Court Decisions
 - Manually annotated categories from Washington University (about 8.4 K files with topic categories)
 - **Eliminate noise:** Regular expressions to identify citations (to other decisions) and the names of legislation.
 - **Different search engine** (Harvard Case Law Access instead of Yahoo)



Experiments

- Data – 8.4K documents
 - 14 “broad” and 279 “narrow” issues (topics)
 - Class size varies from 1 to 1924 documents
- Tested on larger foreground sets
 - Broad issue 1: Criminal Procedures, 1924 cases
 - Broad issue 8: Economic Activity, 1667 cases
 - Broad Issue 5” Privacy, 110 cases
 - Narrow issue 10050: Search and Seizure, 238 cases
 - Subtopic of Broad 1
 - Narrow issue 80010: Antitrust 216 issues
 - Subtopic of Broad 8



Precision for Broad Issue 8

- Baseline system: 23%
- Parameter adjustments: 35%
- Case/legislation filter: 44%
- Additional filter for digits and hyphens: 50%
- Legal Search Customization: 63%
 - Absolute Improvement over baseline: 40%
 - Relative Improvement: 274%



Precision for all 4 tests

Issue	Generality	Freq	Baseline	Final
Criminal 01	Broad	1924	25%	65%
Economic 08	Broad	1667	23%	63%
Privacy 05	Broad	110	27%	40%
Search & Seizure 10050	Narrow	238	19%	30%
Antitrust 80010	Narrow	216	13%	28%

- Narrower & High Frequency Classes have better results
- Domain customization improves results



Background Selection

- **Nordquist and Meyers (2022)**
- Some Terminology systems assume the same background corpus for all foregrounds
 - **We show that this might not be the optimal strategy**
- Cooperative Patent Classification
 - ontology of topic codes for patents
 - Possible to identify classes and superclasses
- For foreground F, experiments with different backgrounds going from the most similar to F (an immediate superclass) to the least similar (a mixture of patents and non patents).
- Example on next 2 slides



Background Sets for SemiConductor Patents

- General = combination of **Base** and OANC
- **Base** = Randomly selected patents
- H = Electricity
- H/01= Basic Electric Elements
- H/01/L = Semiconductor Devices
- H/01/L/21 = **Foreground Patents**: Processes or apparatus adapted for the manufacture or treatment of semiconductor or solid state devices or of parts thereof

IR and Related Applications



Precision of Semiconductor patents with Different Backgrounds

- Results
 - $H/01/L = .63$
 - $H/01 = .61$
 - **$H = .72$**
 - **Base = .7**
 - General = .45
- Best results for a “sweet” spot (the right level of superclass).
- Different types of terminology, depending on the background (e.g., general patent terminology with general corpus, more specific for H/01 then H, etc.)
- Similar results with other topics



Bugs to Fix

- Chinese Termolator
 - Accessor Variety filter fail to produce output files, given the topic: 围棋围棋
- French Termolator
 - Test with: apprentissage automatique
 - Code for French missing in auto_summary.py
 - See “elif language_acronym == ‘fr’
 - See “how run_termolator_fr.sh” should be used



Research Question:

How Can We evaluate Glossaries?

- Some obvious things we can do first:
 - Evaluate sample 100 (or more) for Precision
 - Assume well-formed terms are correct
 - Evaluate sample 100 for Relevance
 - Of the correct terms, what percent is related to foreground topic?
- How do we measure recall for a topic?
 - Find previously created manual glossaries
 - Measure percent coverage
- How do we measure glossary text quality automatically?
 - Find existing online glossaries
 - Compare our glossaries using some existing measure used for Natural Language Generation, e.g., BLEU or ROUGE
- Manual Measures of Glossary Utility
 - Likert Scales
 - Experts or non-experts?
 - What is the key question?
 - Inter-annotator agreement
- After English, how difficult would it be to replicate for French and Chinese?

